

# ENPM673

## PERCEPTION FOR AUTONOMOUS ROBOTS PROJECT-2 REPORT

VENKATA SAI SRICHARAN KASTURI

UID:119444788

## 1Q)

In this problem, you will perform camera pose estimation using homography. Given this video your task is to compute the rotation and translation between the camera and a coordinate frame whose origin is located on any one corner of the sheet of paper.

### Problems encountered:

- 1) Faced issues while producing hough lines.
- 2) Extracting the intersection points(corners), since the multiple points are detected.

### Math:

## Hough transform algorithm



Typically use a different parameterization

$$d = x \cos \theta + y \sin \theta$$

- $d$  is the perpendicular distance from the line to the origin
- $\theta$  is the angle this perpendicular makes with the  $x$  axis
- Why?

### Basic Hough transform algorithm

1. Initialize  $H[d, \theta] = 0$  New discretized image in  $d, \theta$  domain
2. for each edge point  $I[x, y]$  in the image  
    for  $\theta = 0$  to  $180$   
         $d = x \cos \theta + y \sin \theta$   
         $H[d, \theta] += 1$
3. Find the value(s) of  $(d, \theta)$  where  $H[d, \theta]$  is maximum
4. The detected line in the image is given by  $d = x \cos \theta + y \sin \theta$



That gives:


$$\begin{bmatrix} x \\ y \\ 1 \end{bmatrix} = \underbrace{\begin{bmatrix} f & 0 & 0 \\ 0 & f & 0 \\ 0 & 0 & 1 \end{bmatrix}}_{\text{Calibration matrix } K} \begin{bmatrix} X \\ Y \\ Z \end{bmatrix}$$

In short  $\mathbf{x} = K \tilde{\mathbf{X}}$  (here  $\tilde{\mathbf{X}}$  means inhomogeneous coordinates)

**Intrinsic Camera Calibration** means we know  $K$  (we do that later)

We can go from image points into the 3D world:  $\tilde{\mathbf{X}} = K^{-1} \mathbf{x}$

## Adding principal point into $K$




Principal point  $(p_x, p_y)$

Projection with principal point :  $y = f \frac{Y}{Z} + p_y = \frac{fY + Zp_y}{Z}$  and  $x = f \frac{X}{Z} + p_x = \frac{fX + Zp_x}{Z}$

That gives:

$$\begin{bmatrix} x \\ y \\ 1 \end{bmatrix} = \begin{bmatrix} f & 0 & p_x \\ 0 & f & p_y \\ 0 & 0 & 1 \end{bmatrix} \begin{bmatrix} X \\ Y \\ Z \end{bmatrix}$$



Given a 3D homogenous point  $X^w$  in world coordinate system

- 1) Translate from world to camera coordinate system:

$$\tilde{X}^{c'} = \tilde{X}^w - \tilde{C}$$

$$\tilde{X}^{c'} = \underbrace{(I_{3 \times 3} \mid -\tilde{C})}_{3 \times 4 \text{ matrix}} X^w \quad \text{where } I_{3 \times 3} \text{ is } 3 \times 3 \text{ identity matrix}$$

Last column is translation

- 2) Rotate world coordinate system into camera coordinate system

$$\tilde{X}^c = R (I_{3 \times 3} \mid -\tilde{C}) X^w$$

- 3) Apply camera matrix

$$x = K R (I_{3 \times 3} \mid -\tilde{C}) X$$

Tilda means inhomogeneous vector

## Solving for Homographies

$$\begin{bmatrix} x'_i \\ y'_i \\ 1 \end{bmatrix} \cong \begin{bmatrix} h_{00} & h_{01} & h_{02} \\ h_{10} & h_{11} & h_{12} \\ h_{20} & h_{21} & h_{22} \end{bmatrix} \begin{bmatrix} x_i \\ y_i \\ 1 \end{bmatrix}$$

Camera coordinates

$$\begin{aligned} x'_i &= \frac{h_{00}x_i + h_{01}y_i + h_{02}}{h_{20}x_i + h_{21}y_i + h_{22}} \\ y'_i &= \frac{h_{10}x_i + h_{11}y_i + h_{12}}{h_{20}x_i + h_{21}y_i + h_{22}} \end{aligned}$$

$$x'_i(h_{20}x_i + h_{21}y_i + h_{22}) = h_{00}x_i + h_{01}y_i + h_{02}$$

$$y'_i(h_{20}x_i + h_{21}y_i + h_{22}) = h_{10}x_i + h_{11}y_i + h_{12}$$

$$\begin{bmatrix} x_i & y_i & 1 & 0 & 0 & 0 & -x'_i x_i & -x'_i y_i & -x'_i \\ 0 & 0 & 0 & x_i & y_i & 1 & -y'_i x_i & -y'_i y_i & -y'_i \end{bmatrix} \begin{bmatrix} h_{00} \\ h_{01} \\ h_{02} \\ h_{10} \\ h_{11} \\ h_{12} \\ h_{20} \\ h_{21} \\ h_{22} \end{bmatrix} = \begin{bmatrix} 0 \\ 0 \end{bmatrix}$$

# Camera pose estimation

Assume all points lie in one plane with  $Z=0$ :

$$X = (X, Y, 0, 1)$$

$$x = PX$$

$$= K[r_1 r_2 r_3 t] \begin{pmatrix} X \\ Y \\ 0 \\ 1 \end{pmatrix}$$

$$= K[r_1 r_2 t] \begin{pmatrix} X \\ Y \\ 1 \end{pmatrix}$$

$$= H \begin{pmatrix} X \\ Y \\ 1 \end{pmatrix}$$

$$H = \lambda K[r_1 r_2 t]$$

$$K^{-1}H = \lambda[r_1 r_2 t]$$

– $r_1$  and  $r_2$  are unit vectors  $\Rightarrow$  find lambda

–Use this to compute t


–Rotation matrices are orthogonal  $\Rightarrow$  find  $r_3$

$$P = K \begin{bmatrix} r_1 & r_2 & (r_1 \times r_2) & t \end{bmatrix}$$

## Steps to be followed:

- 1) Open video file and read frame by frame.
- 2) Apply blurring to the current frame using a 10x10 kernel.
- 3) Convert the blurred frame from BGR color space to HSV color space.
- 4) Define a range of white color and create a mask using cv.inRange().
- 5) Detect edges in the masked image using cv.Canny() function.
- 6) Apply Hough transform to find the lines in the edge image using the hough\_lines() function.
- 7) Find the four intersection points of the lines detected in the previous step using get\_intersection() function.
- 8) Store these points as corners of the object in the image.

- 9) Use the corners of the object to calculate homography using `find_homography()` function.
- 10) Use the homography matrix to calculate the pose of the object in the image using `find_rotation_translation()` function.
- 11) Extract the roll, pitch, and yaw values from the pose matrix and store them in the respective lists.
- 12) Extract the x, y, and z values from the pose matrix and store them in the respective lists.
- 13) Display the frame with the detected lines and intersection points.
- 14) Continue the loop until all frames have been processed.
- 15) Release the video object and close all windows.



## Pipeline

## Code:

```
import numpy as np
from matplotlib import pyplot as plt
import cv2 as cv,os
import math
CURRENT_DIR = os.path.dirname(__file__)
path_video=os.path.join(CURRENT_DIR,'project2.avi')
video= cv.VideoCapture(path_video)

coordinates = np.array([[0, 0], [0, 27.9], [21.6,27.9], [21.6, 0]])

def find_homography(p1, p2):

    if p1.shape[0] < 4 or p2.shape[0] < 4:
        raise ValueError("Not enough points to find homography")

    A = []
    for i in range(p1.shape[0]):
        x, y = p1[i]
        u, v = p2[i]
        A.append([-x, -y, -1, 0, 0, 0, x*u, y*u, u])
        A.append([0, 0, 0, -x, -y, 1, x*v, y*v, v])
    A = np.array(A)

    # SVD
    _, _, Vt = np.linalg.svd(A)
    H = Vt[-1].reshape(3, 3)
    H /= H[2, 2]

    return H

def find_rotation_translation(H, K):

    K_inv = np.linalg.inv(K)
    # Vectors
    r1 = H[:, 0:1] / np.linalg.norm(H[:, 0])
    r2 = H[:, 1:2] / np.linalg.norm(H[:, 1])
    r3 = np.cross(r1.T, r2.T).T

    # Rotational matrix
    R = np.hstack([r1,r2,r3])

    # Translation
```



```

    t = K_inv @ H[:, 2] / np.linalg.norm(K_inv @ H[:, 2])

# Pose
    P = np.hstack([R, t.reshape(-1, 1)])

    return P

def hough_lines(edges):

    theta = np.arange(-90, 90)
    cos_t = np.cos(theta)
    sin_t = np.sin(theta)
    a=edges.shape[0]
    b=edges.shape[1]
    c=int(np.ceil(np.sqrt(a**2+b**2)))
    perpendicular_range=np.arange(-c,c+1,1)
    accumulator = np.zeros((len(perpendicular_range), len(theta)))
    y_indexes, x_indexes = np.nonzero(edges)

    for i in range(len(x_indexes)):
        x = x_indexes[i]
        y = y_indexes[i]

        for j in range(len(theta)):
            rho = int(round(x * cos_t[j] + y * sin_t[j])) + c
            accumulator[rho, j] += 1

# Maximum votes
    rhos, thetas = np.nonzero(accumulator >=55)

# Conversion
    lines = []
    for i in range(len(rhos)):
        rho = perpendicular_range[rhos[i]]
        angle = theta[thetas[i]]
        lines.append((rho, angle))

# Polar to cartesian
    cartesian_lines = []
    for line in lines:
        rho, theta = line
        a = np.cos(theta)
        b = np.sin(theta)
        x0 = a * (rho)
        y0 = b * (rho)
        x1 = int(x0 + 1500 * (-b))

```

```

        y1 = int(y0 + 1500 * (a))
        x2 = int(x0 - 1500 * (-b))
        y2 = int(y0 - 1500 * (a))
        cartesian_lines.append([(x1, y1), (x2, y2)])
        cv.line(frame, (x1, y1), (x2, y2), (255, 0, 0), 1)
# Intersection points
    intersections = []
    for i in range(len(cartesian_lines)):
        for j in range(i+1, len(cartesian_lines)):
            line1 = cartesian_lines[i]
            line2 = cartesian_lines[j]
            intersection = get_intersection(line1, line2)
            if intersection is not None:
                intersections.append(intersection)

    intersections.sort()
    corners = np.array(intersections)

    return corners[:4, :]

def get_intersection(line1, line2):

    x1, y1 = line1[0]
    x2, y2 = line1[1]
    x3, y3 = line2[0]
    x4, y4 = line2[1]

    denominator = (x1-x2)*(y3-y4) - (y1-y2)*(x3-x4)
    if denominator == 0:

        return None
    else:
        px = ((x1*y2-y1*x2)*(x3-x4)-(x1-x2)*(x3*y4-y3*x4))/denominator
        py = ((x1*y2-y1*x2)*(y3-y4)-(y1-y2)*(x3*y4-y3*x4))/denominator
        return (int(px), int(py))

roll_list, pitch_list, yaw_list = [], [], []
x_list, y_list, z_list = [], [], []

while(video.isOpened()):

    ret, frame=video.read()
    if ret==True:

```

```

        blur = cv.blur(frame,(10,10))
        hsv=cv.cvtColor(blur,cv.COLOR_BGR2HSV)
        whitepaper1=np.array([0,0,235])
        whitepaper2=np.array([180,100,255])
        mask=cv.inRange(hsv,whitepaper1,whitepaper2)
        edge=cv.Canny(mask,10,200)
        g=hough_lines(edge)
        T=find_rotation_translation(find_homography(g,coordinates),K =
np.array([[1.38e+03, 0, 9.46e+02], [0, 1.38e+03, 5.27e+02], [0, 0, 1]]))
        R=T[:, :3]
        t=T[:, 3]

        roll, pitch, yaw = [math.atan2(R[i, j], R[i, k]) for i, j, k in [(0, 1,
2), (1, 2, 0), (2, 0, 1)]]
        tx, ty, tz = t

        roll_list.append(roll)
        pitch_list.append(pitch)
        yaw_list.append(yaw)
        x_list.append(tx)
        y_list.append(ty)
        z_list.append(tz)

        cv.imshow('lines',frame)
        # cv.imshow('frame',edge)
        key =cv.waitKey(1)
        if key==ord('q'):
            break
    else:
        break
video.release()
cv.destroyAllWindows()


plt.plot(roll_list, label='Roll',color='Teal')
plt.plot(pitch_list, label='Pitch',color='Gold')
plt.plot(yaw_list, label='Yaw',color='Red')
plt.legend()
plt.xlabel('Frame')
plt.ylabel('Angle (rad)')
plt.show()



plt.plot(x_list, label='X',color='Blue')
plt.plot(y_list, label='Y',color='Orange')
plt.plot(z_list, label='Z',color='Green')
plt.legend()

```

```
plt.xlabel('Frame')  
plt.ylabel('Translation (cm)')  
plt.show()
```

**Output:**





## Q2)

You are given four images which were taken from the same camera position (only rotation no translation) you will need to stitch these images to create a panoramic image.

### Problems encountered:

- 1) Stitching the images. Blank spaces occurred for the final panorama.
- 2) Pano1 in the final result is not being displayed correctly. (Note: Using the built in homography function the result is good)

### Math:

## Solving for Homographies

$$\begin{bmatrix} x'_i \\ y'_i \\ 1 \end{bmatrix} \cong \begin{bmatrix} h_{00} & h_{01} & h_{02} \\ h_{10} & h_{11} & h_{12} \\ h_{20} & h_{21} & h_{22} \end{bmatrix} \begin{bmatrix} x_i \\ y_i \\ 1 \end{bmatrix}$$

Camera coordinates


→

$$\begin{aligned} x'_i &= \frac{h_{00}x_i + h_{01}y_i + h_{02}}{h_{20}x_i + h_{21}y_i + h_{22}} \\ y'_i &= \frac{h_{10}x_i + h_{11}y_i + h_{12}}{h_{20}x_i + h_{21}y_i + h_{22}} \end{aligned}$$

$$\begin{aligned} x'_i(h_{20}x_i + h_{21}y_i + h_{22}) &= h_{00}x_i + h_{01}y_i + h_{02} \\ y'_i(h_{20}x_i + h_{21}y_i + h_{22}) &= h_{10}x_i + h_{11}y_i + h_{12} \end{aligned}$$
$$\begin{bmatrix} x_i & y_i & 1 & 0 & 0 & 0 & -x'_i x_i & -x'_i y_i & -x'_i \\ 0 & 0 & 0 & x_i & y_i & 1 & -y'_i x_i & -y'_i y_i & -y'_i \end{bmatrix} \begin{bmatrix} h_{00} \\ h_{01} \\ h_{02} \\ h_{10} \\ h_{11} \\ h_{12} \\ h_{20} \\ h_{21} \\ h_{22} \end{bmatrix} = \begin{bmatrix} 0 \\ 0 \end{bmatrix}$$

## **Steps to be followed:**

- 1) Import necessary libraries: Begin by importing the necessary libraries, such as OpenCV, NumPy, and Matplotlib. This will allow you to perform the required operations.
- 2) Load the images: Load the images that you want to stitch together. In this case, four images are loaded and stored in the `image_paths` list.
- 3) Resize the images: Resize the images to ensure that they have a uniform size. In this code, the images are resized to a maximum dimension of 500 pixels while maintaining the aspect ratio.
- 4) Detect keypoints and descriptors: Use SIFT (Scale-Invariant Feature Transform) to detect keypoints and compute descriptors for each image.
- 5) Match the keypoints: Use FLANN (Fast Library for Approximate Nearest Neighbors) to match the keypoints between the images.
- 6) Draw matches: Draw the matched keypoints between the images using the `drawMatchesKnn` function.
- 7) Compute homography: Compute the homography matrix for each pair of images using the matches between them.
- 8) Compute the size of the output image: Use the homography matrix to compute the size of the output image.
- 9) Warp the images: Warp the images using the computed homography matrix.
- 10) Blend the images: Blend the warped images together to create a seamless panorama.
- 11) Display the final panorama: Display the final panorama using Matplotlib.
- 12) Save the output: Save the final panorama to a file for later use.



## Pipeline



## Code:

```
import cv2 as cv, numpy as np, os
import matplotlib.pyplot as plt
sift = cv.SIFT_create()

CURRENT_DIR = os.path.dirname(__file__)

# Importing images
image_paths=[]

for i in range(1,5):
    image_paths.append(os.path.join(CURRENT_DIR, f'image_{i}.jpg'))

img1 = cv.imread(image_paths[0], cv.IMREAD_GRAYSCALE)
h, w = img1.shape[:2]
ratio = float(500) / max(h, w)
new_h = int(ratio * h)
new_w = int(ratio * w)

# Resize the image
img1 = cv.resize(img1, (new_w, new_h))

img2 = cv.imread(image_paths[1], cv.IMREAD_GRAYSCALE)
# Resize the image
img2 = cv.resize(img2, (new_w, new_h))

img3 = cv.imread(image_paths[2], cv.IMREAD_GRAYSCALE)
# Resize the image
img3 = cv.resize(img3, (new_w, new_h))

img4 = cv.imread(image_paths[3], cv.IMREAD_GRAYSCALE)
# Resize the image
img4 = cv.resize(img4, (new_w, new_h))

def trimmer(frame):
    while not np.any(frame[0]):
        frame = frame[1:]
    while not np.any(frame[-1]):
        frame = frame[:-2]
    while not np.any(frame[:,0]):
        frame = frame[:,1:]
    while not np.any(frame[:, -1]):
        frame = frame[:, :-2]
    return frame
```

```

kp1, des1 = sift.detectAndCompute(img1, None)
kp2, des2 = sift.detectAndCompute(img2, None)
kp3, des3 = sift.detectAndCompute(img3, None)
kp4, des4 = sift.detectAndCompute(img4, None)

FLANN_INDEX_KDTREE = 0
index_params = dict(algorithm=FLANN_INDEX_KDTREE, trees=5)
search_params = dict(checks=50)

flann = cv.FlannBasedMatcher(index_params, search_params)

matches12 = flann.knnMatch(des1, des2, k=2)
matches23 = flann.knnMatch(des2, des3, k=2)
matches34 = flann.knnMatch(des3, des4, k=2)

good12 = []
if len(matches12) > 0:
    for m, n in matches12:
        if m.distance < 0.7 * n.distance:
            good12.append([m])

good23 = []
if len(matches23) > 0:
    for m, n in matches23:
        if m.distance < 0.7 * n.distance:
            good23.append([m])

good34 = []
if len(matches34) > 0:
    for m, n in matches34:
        if m.distance < 0.7 * n.distance:
            good34.append([m])

img_matches12 = cv.drawMatchesKnn(img1, kp1, img2, kp2, good12, None,
flags=cv.DrawMatchesFlags_NOT_DRAW_SINGLE_POINTS)
img_matches23 = cv.drawMatchesKnn(img2, kp2, img3, kp3, good23, None,
flags=cv.DrawMatchesFlags_NOT_DRAW_SINGLE_POINTS)
img_matches34 = cv.drawMatchesKnn(img3, kp3, img4, kp4, good34, None,
flags=cv.DrawMatchesFlags_NOT_DRAW_SINGLE_POINTS)

cv.imshow('Matches 1-2', img_matches12)
cv.imshow('Matches 2-3', img_matches23)
cv.imshow('Matches 3-4', img_matches34)
cv.waitKey(0)

```

```

MIN_MATCH_COUNT = 10

def compute_homography(kp1, kp2, matches):
    if len(matches) < MIN_MATCH_COUNT:
        return None

    src_pts = np.float32([ kp1[m[0].queryIdx].pt for m in matches ]).reshape(-1,1,2)
    dst_pts = np.float32([ kp2[m[0].trainIdx].pt for m in matches ]).reshape(-1,1,2)

    src_pts = np.hstack((src_pts.reshape(-1,2), np.ones((len(src_pts), 1))))
    dst_pts = np.hstack((dst_pts.reshape(-1,2), np.ones((len(dst_pts), 1))))

    if src_pts.shape[0] < 4 or dst_pts.shape[0] < 4:
        raise ValueError("Not enough points to find homography")

    A = []
    for i in range(len(src_pts)):
        x, y = src_pts[i,0:2]
        u, v = dst_pts[i,0:2]
        A.append([x,y,1, 0, 0, 0, -x*u, -y*u, -u])
        A.append([0, 0, 0, x*0.5, y, 1, -x*v, -y*v, -v])
    A = np.array(A)

    __, __, Vt = np.linalg.svd(A)
    H = Vt[-1].reshape(3, 3)
    H /= H[2, 2]

    return H

H12 = compute_homography(kp1, kp2, good12)
H23 = compute_homography(kp2, kp3, good23)
H34 = compute_homography(kp3, kp4, good34)

# Compute the size of the output image
corners1 = np.float32([[0, 0], [0, h], [w, h], [w, 0]]).reshape(-1, 1, 2)
corners2 = cv.perspectiveTransform(corners1, H12)

```

```

corners3 = cv.perspectiveTransform(corners2, H23)
corners4 = cv.perspectiveTransform(corners3, H34)

x_min = int(min(corners1[:, :, 0].min(), corners2[:, :, 0].min(),
corners3[:, :, 0].min(), corners4[:, :, 0].min()))
y_min = int(min(corners1[:, :, 1].min(), corners2[:, :, 1].min(),
corners3[:, :, 1].min(), corners4[:, :, 1].min()))
x_max = int(max(corners1[:, :, 0].max(), corners2[:, :, 0].max(),
corners3[:, :, 0].max(), corners4[:, :, 0].max()))
y_max = int(max(corners1[:, :, 1].max(), corners2[:, :, 1].max(),
corners3[:, :, 1].max(), corners4[:, :, 1].max()))

# Create a translation matrix to shift the image
translation_matrix = np.array([[1, 0, -x_min], [0, 1, -y_min], [0, 0, 1]])

# Warp the images using the homographies
pano1 = cv.warpPerspective(img1, translation_matrix.dot(H12), (x_max - x_min,
y_max - y_min))
pano1=trimmer(pano1)
# pano1=cv.cvtColor(pano1,cv.COLOR_GRAY2BGR)
pano2 = cv.warpPerspective(img2, translation_matrix.dot(H23), (x_max - x_min,
y_max - y_min))
pano2=trimmer(pano2)
pano3 = cv.warpPerspective(img3, translation_matrix.dot(H34), (x_max - x_min,
y_max - y_min))
pano3=trimmer(pano3)

cv.imshow('Panorama1', pano1)
cv.imshow('Panorama2', pano2)
cv.imshow('Panorama3', pano3)
cv.waitKey(0)

print(pano1.shape, pano2.shape, pano3.shape)
# Combine the images
pano =
np.zeros((pano1.shape[0]+pano2.shape[0]+pano3.shape[0], pano1.shape[1]+pano2.shape
[1]+pano3.shape[1]), dtype=np.uint8)
# print(pano.shape)
pano[:pano1.shape[0], :pano1.shape[1]] = cv.flip(pano1, 0)
pano[:pano1.shape[0], :pano1.shape[1]] = pano1
# print(pano2.shape, pano1.shape[1]+pano2.shape[1])
pano[:pano3.shape[0], pano1.shape[1]+pano2.shape[1]:] = cv.flip(pano3, 0)
pano[:pano3.shape[0], pano1.shape[1]+pano2.shape[1]:] = pano3
pano[:pano2.shape[0], pano1.shape[1]:pano1.shape[1]+pano2.shape[1]] =
cv.flip(pano2, 0)

```

```
pano[:pano2.shape[0], pano1.shape[1]:pano1.shape[1]+pano2.shape[1]] = pano2

# Show the resulting panorama
pano=trimmer(pano)
print(pano.shape,'shape')
pano=pano[100:550,:]
cv.imshow('Panorama', pano)

cv.waitKey(0)
```

**Output:**



**Matches 1-2**





**Matches2-3**



**Matches3-4**



**Final panorama**